

Lecture 8: Camera Models

Dr. Juan Carlos Niebles
Stanford AI Lab

Professor Fei-Fei Li
Stanford Vision Lab

What we will learn today?

- Pinhole cameras
- Cameras & lenses
- The geometry of pinhole cameras
 - Projection matrix
 - Intrinsic parameters
 - Extrinsic parameters

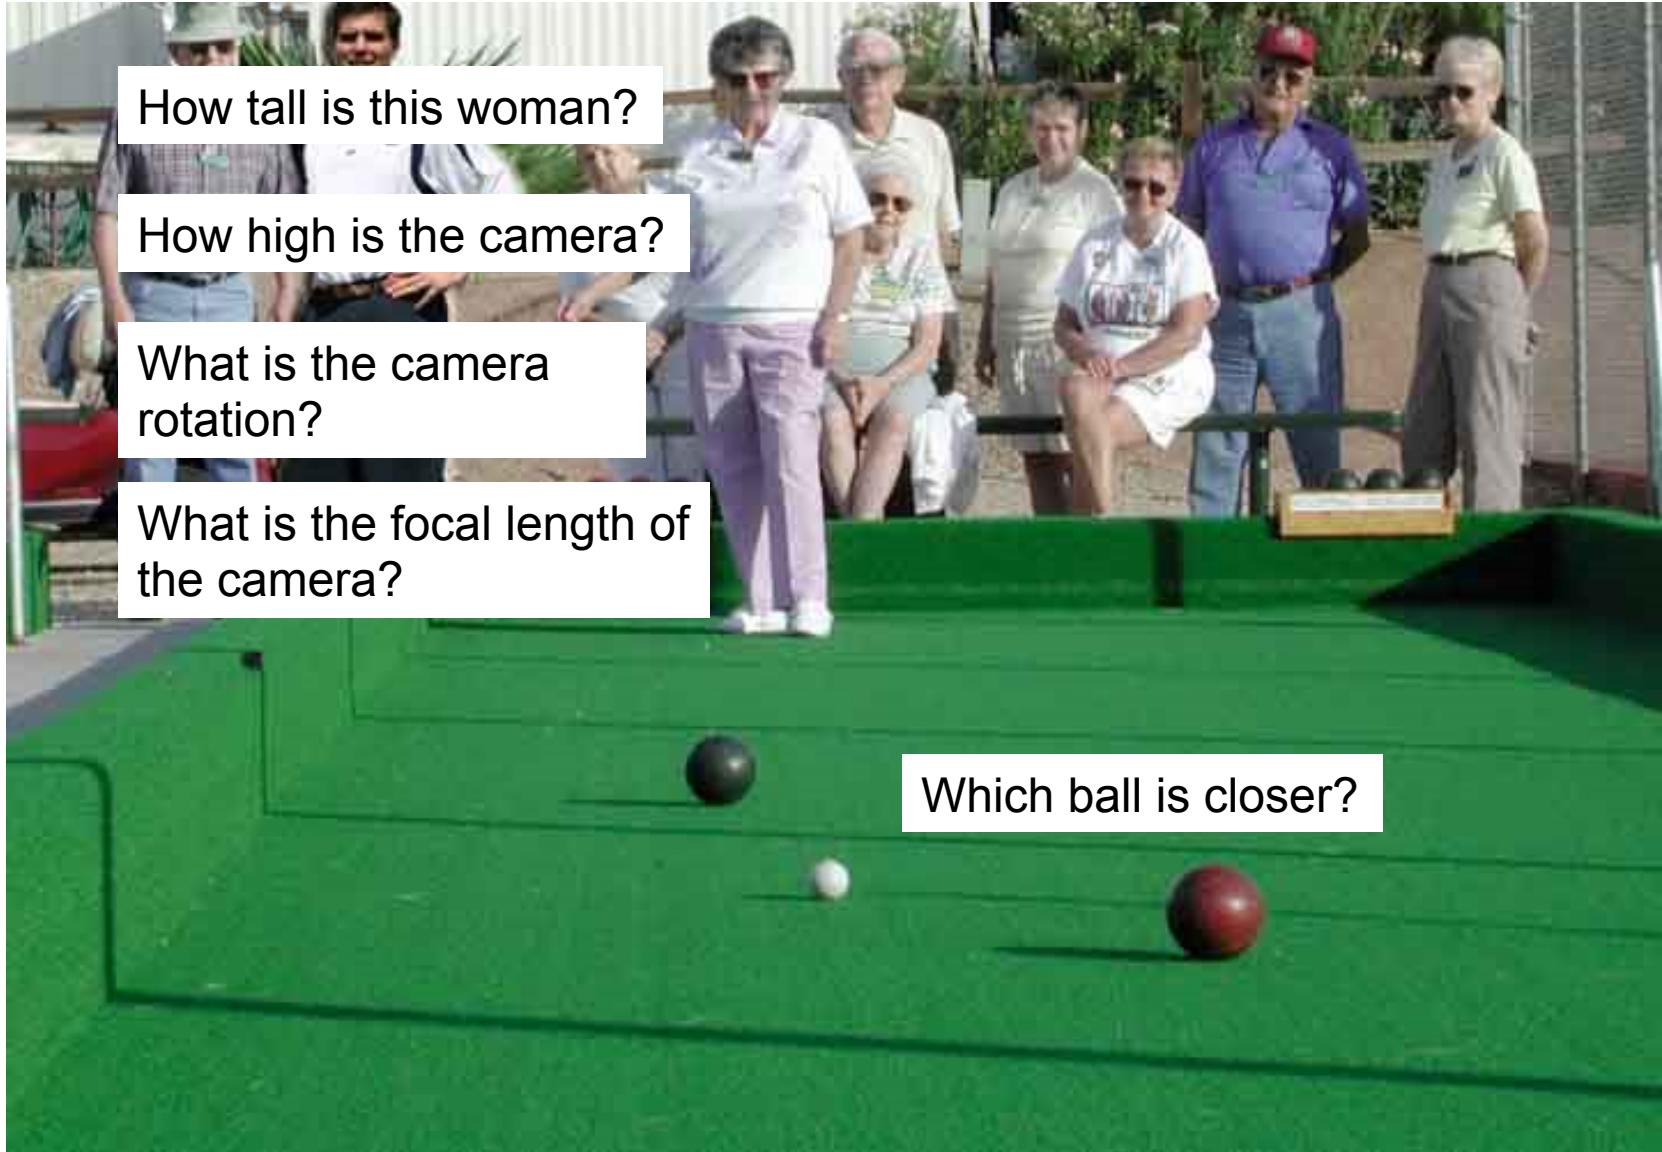
Reading:
[FP] Chapters 1 – 3
[HZ] Chapter 6

What we will learn today?

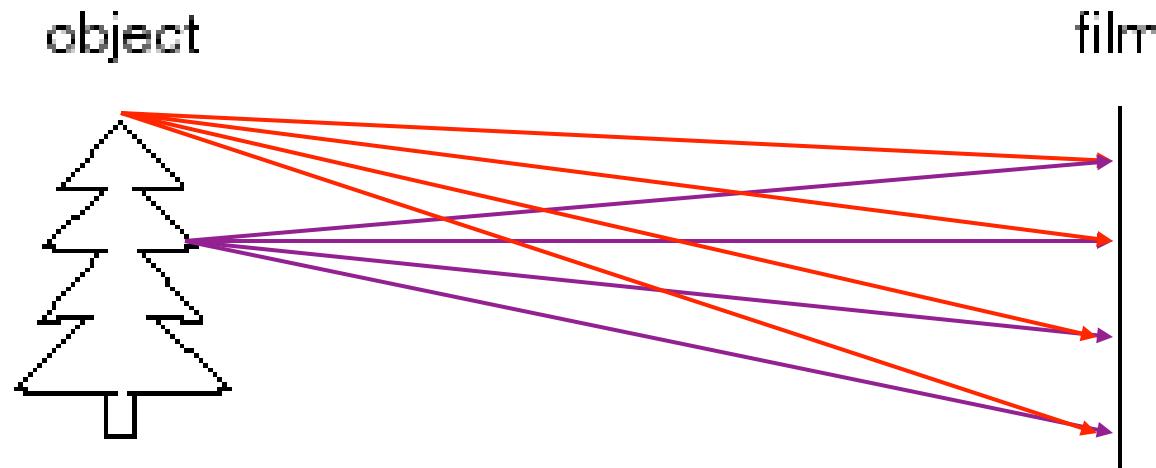
- Pinhole cameras
- Cameras & lenses
- The geometry of pinhole cameras
 - Projection matrix
 - Intrinsic parameters
 - Extrinsic parameters

Reading:
[FP] Chapters 1 – 3
[HZ] Chapter 6

Camera and World Geometry

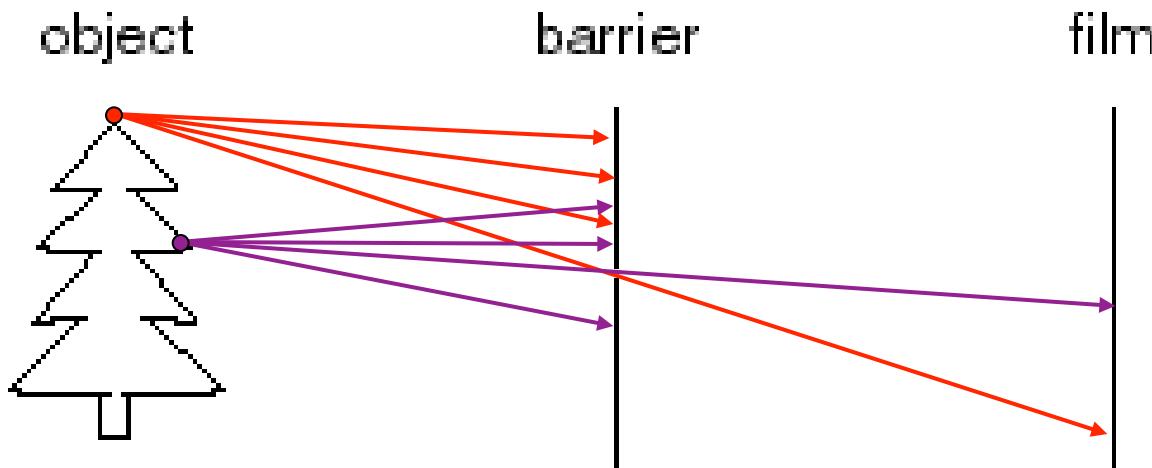


How do we see the world?



- Let's design a camera
 - Idea 1: put a piece of film in front of an object
 - Do we get a reasonable image?

Pinhole camera



- Add a barrier to block off most of the rays
 - This reduces blurring
 - The opening known as the **aperture**

Camera obscura: the pre-camera

- Known during classical period in China and Greece
(e.g. Mo-Tsi, China, 470BC to 390BC)

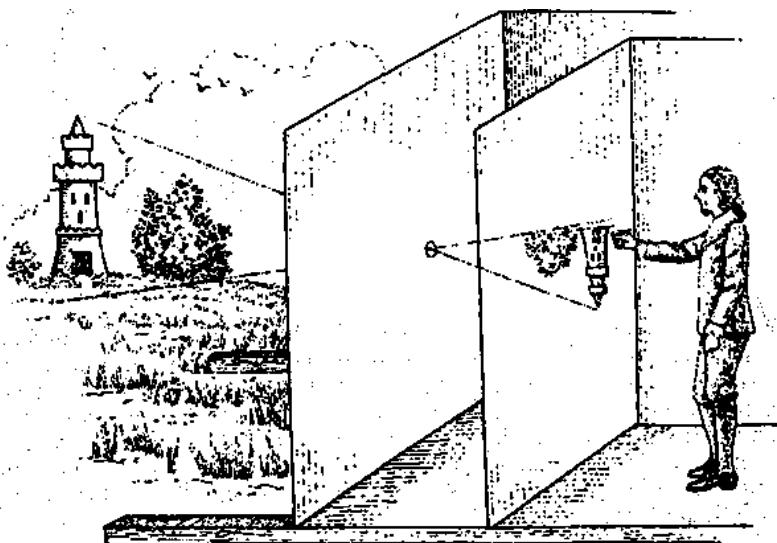


Illustration of Camera Obscura

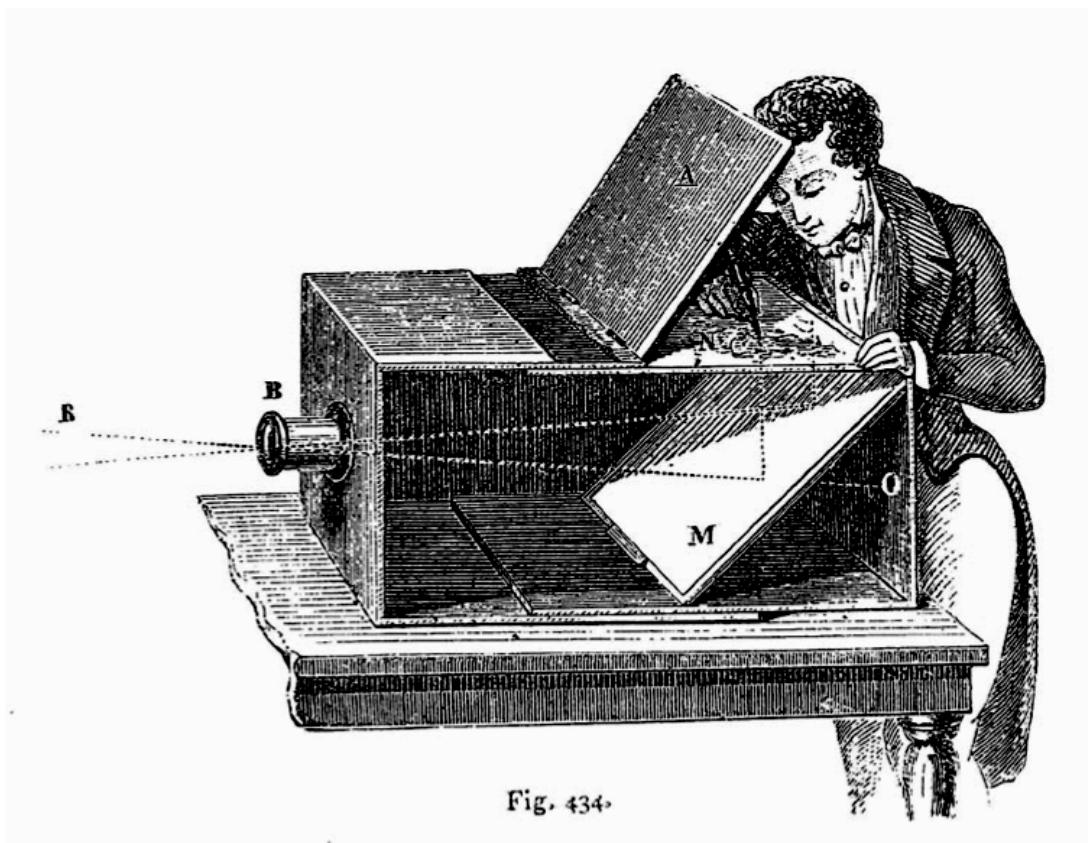


Freestanding camera obscura at UNC Chapel Hill

Photo by Seth Ilys

Slide credit: J. Hayes

Camera Obscura used for Tracing



Lens Based Camera Obscura, 1568

Slide credit: J. Hayes

First Photograph

Oldest surviving photograph

- Took 8 hours on pewter plate



Joseph Niepce, 1826

Photograph of the first photograph




Stored at UT Austin

Niepce later teamed up with Daguerre, who eventually created Daguerrotypes

Slide credit: J. Hayes

Dimensionality Reduction Machine (3D to 2D)



Figures © Stephen E. Palmer, 2002

Projection can be tricky...



Slide source: Seitz

Projection can be tricky...

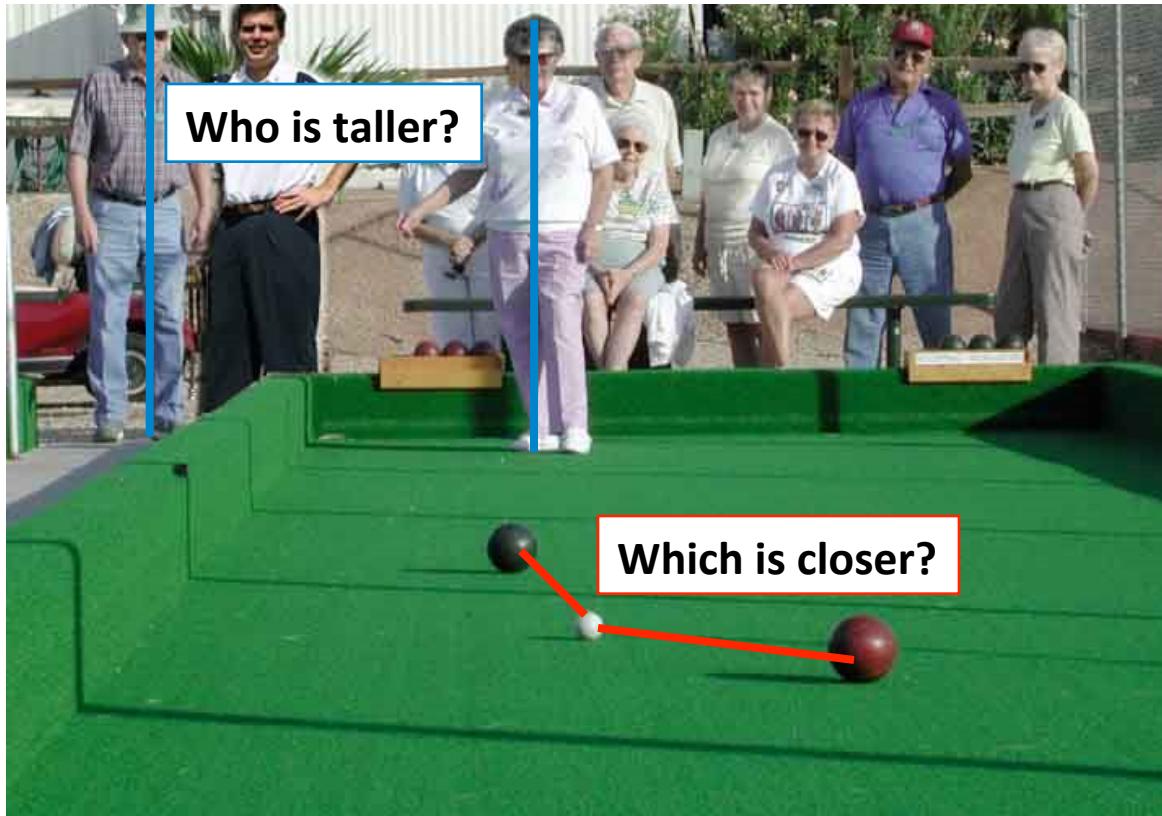


Slide source: Seitz

Projective Geometry

What is lost?

- Length



Slide credit: J. Hayes

Length is not preserved

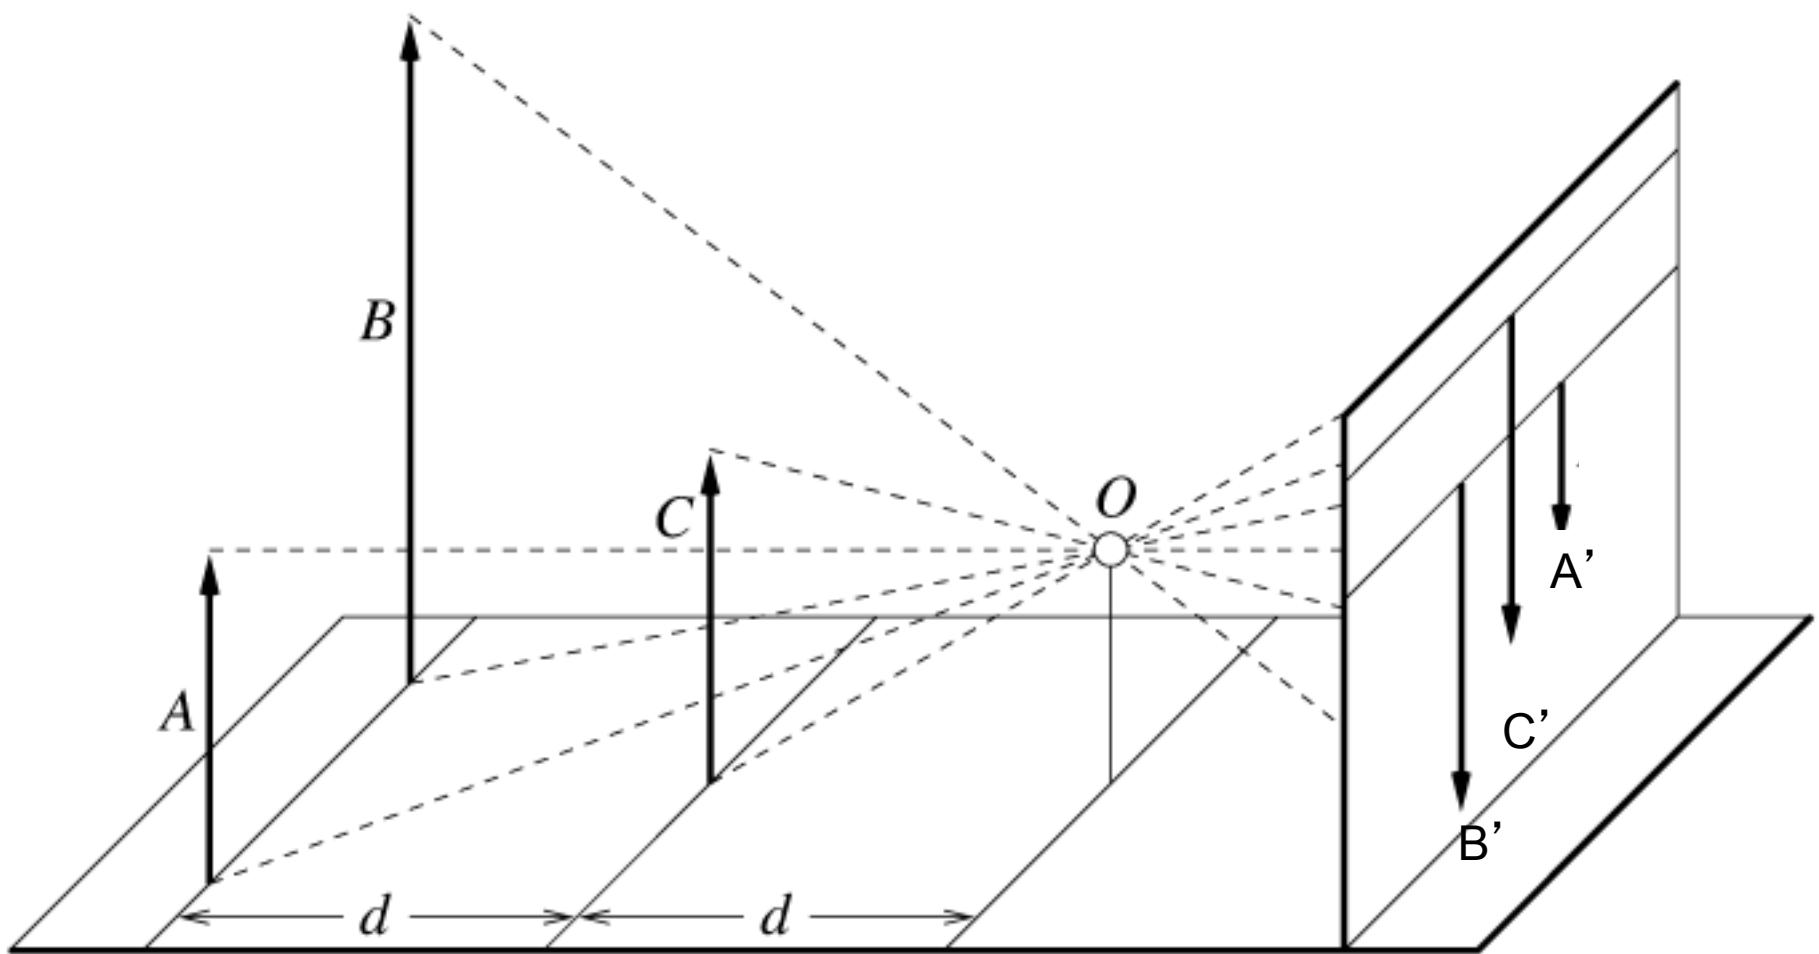
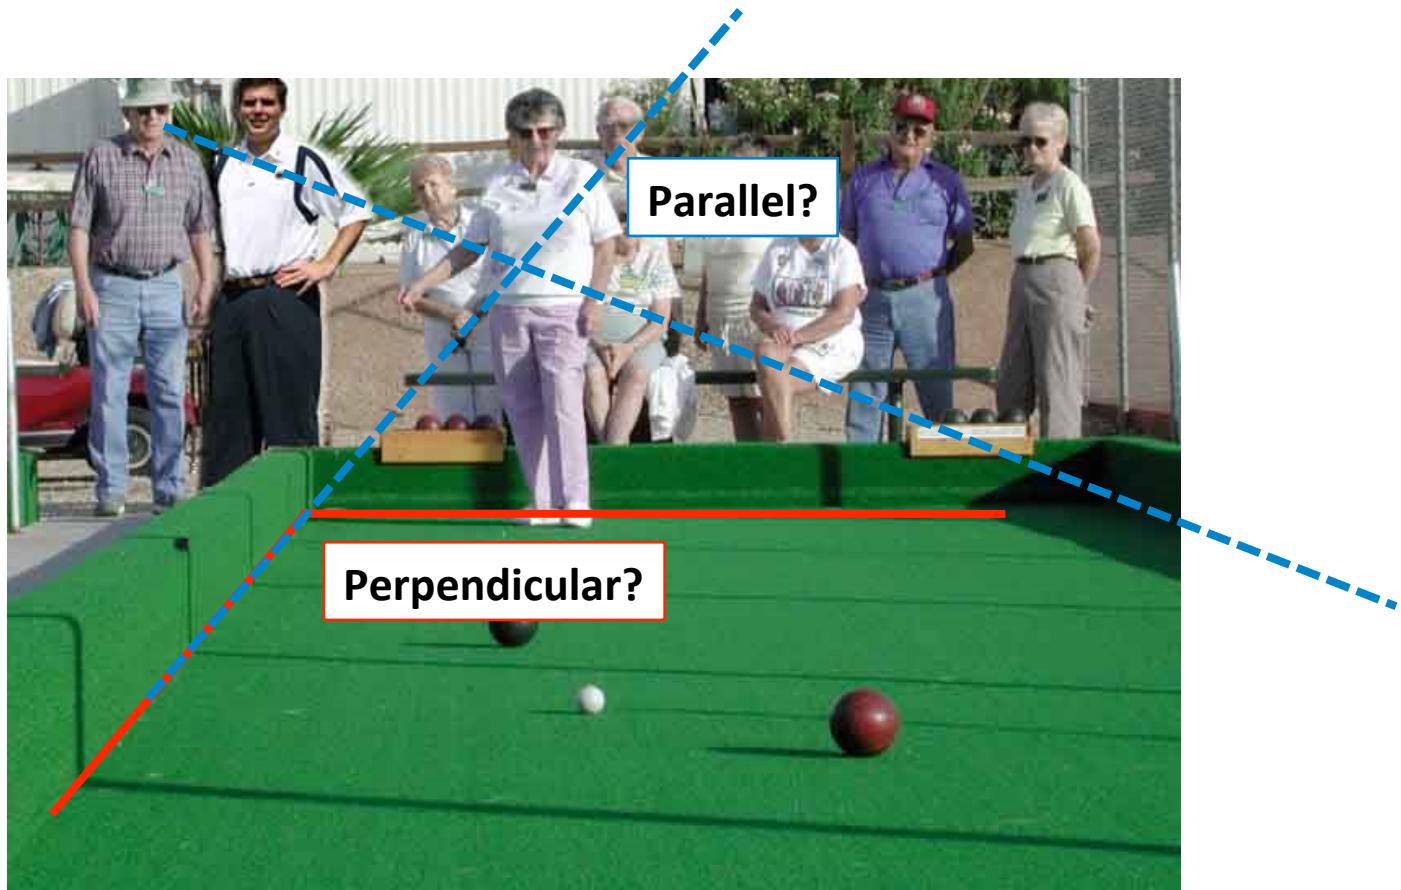


Figure by David Forsyth

Projective Geometry

What is lost?

- Length
- Angles

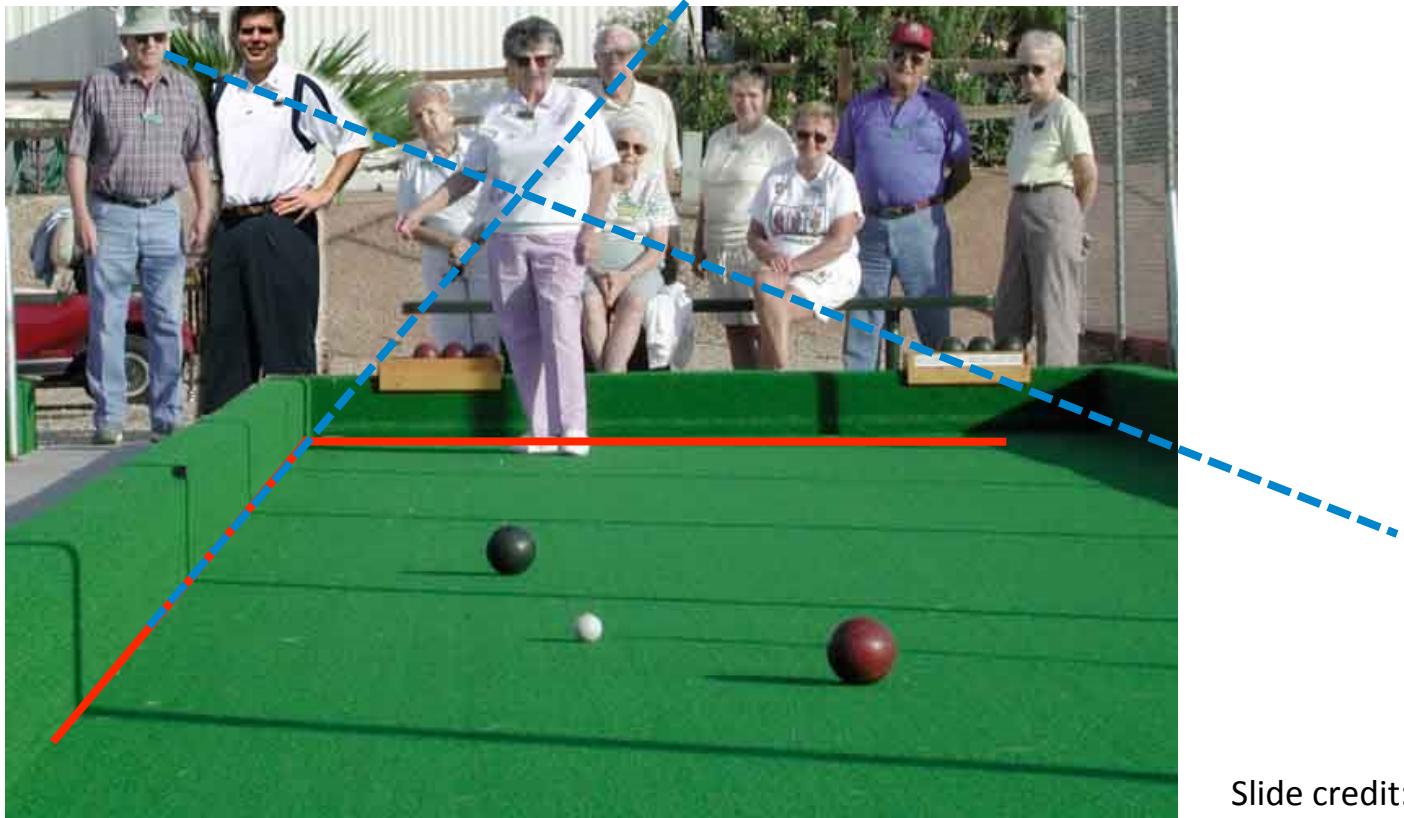


Slide credit: J. Hayes

Projective Geometry

What is preserved?

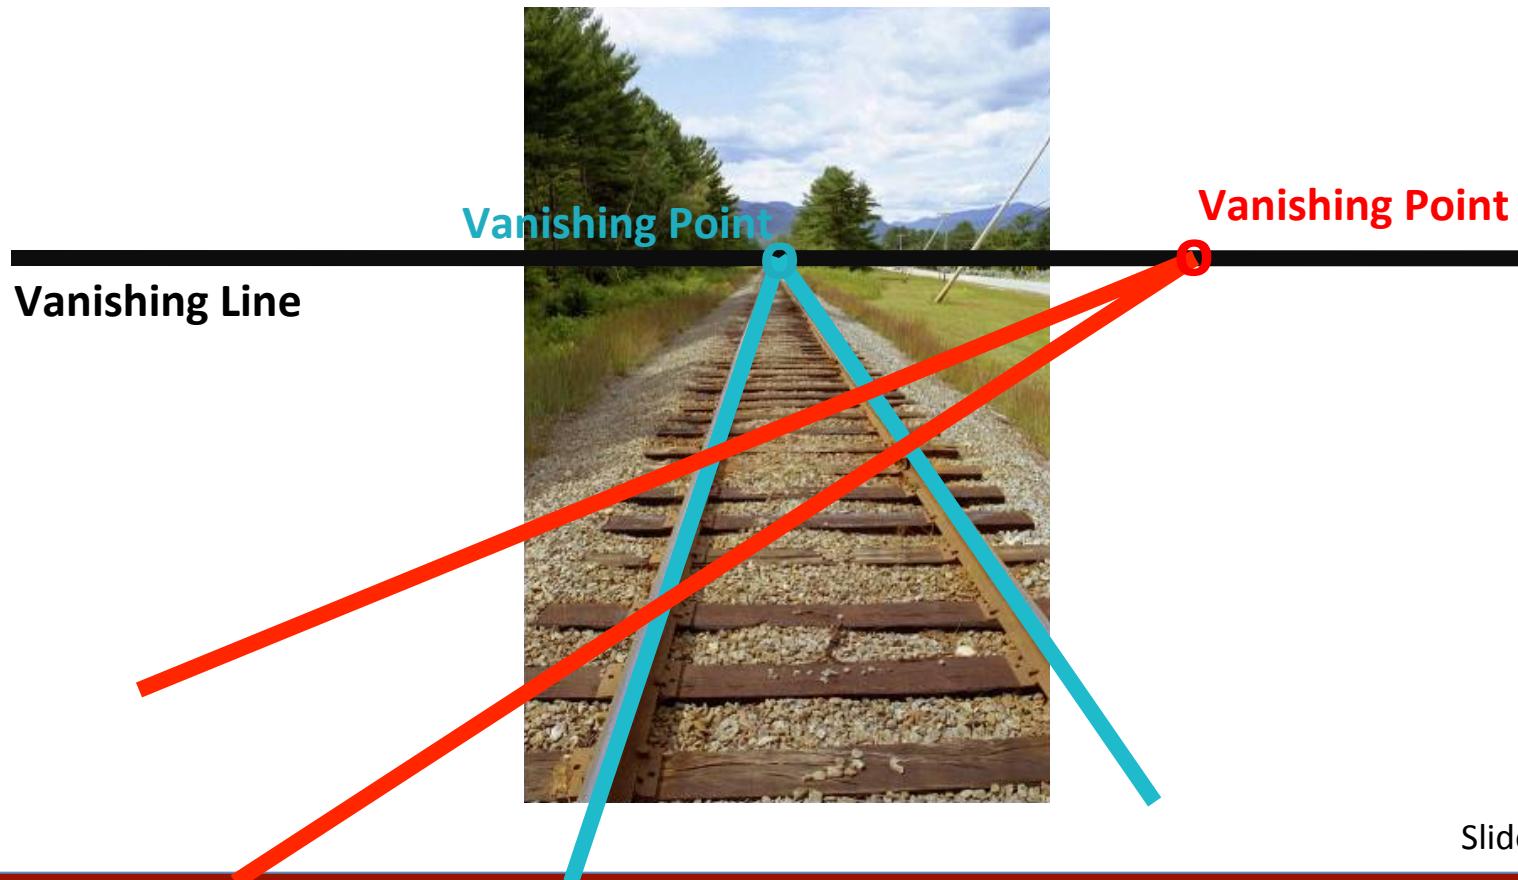
- Straight lines are still straight



Slide credit: J. Hayes

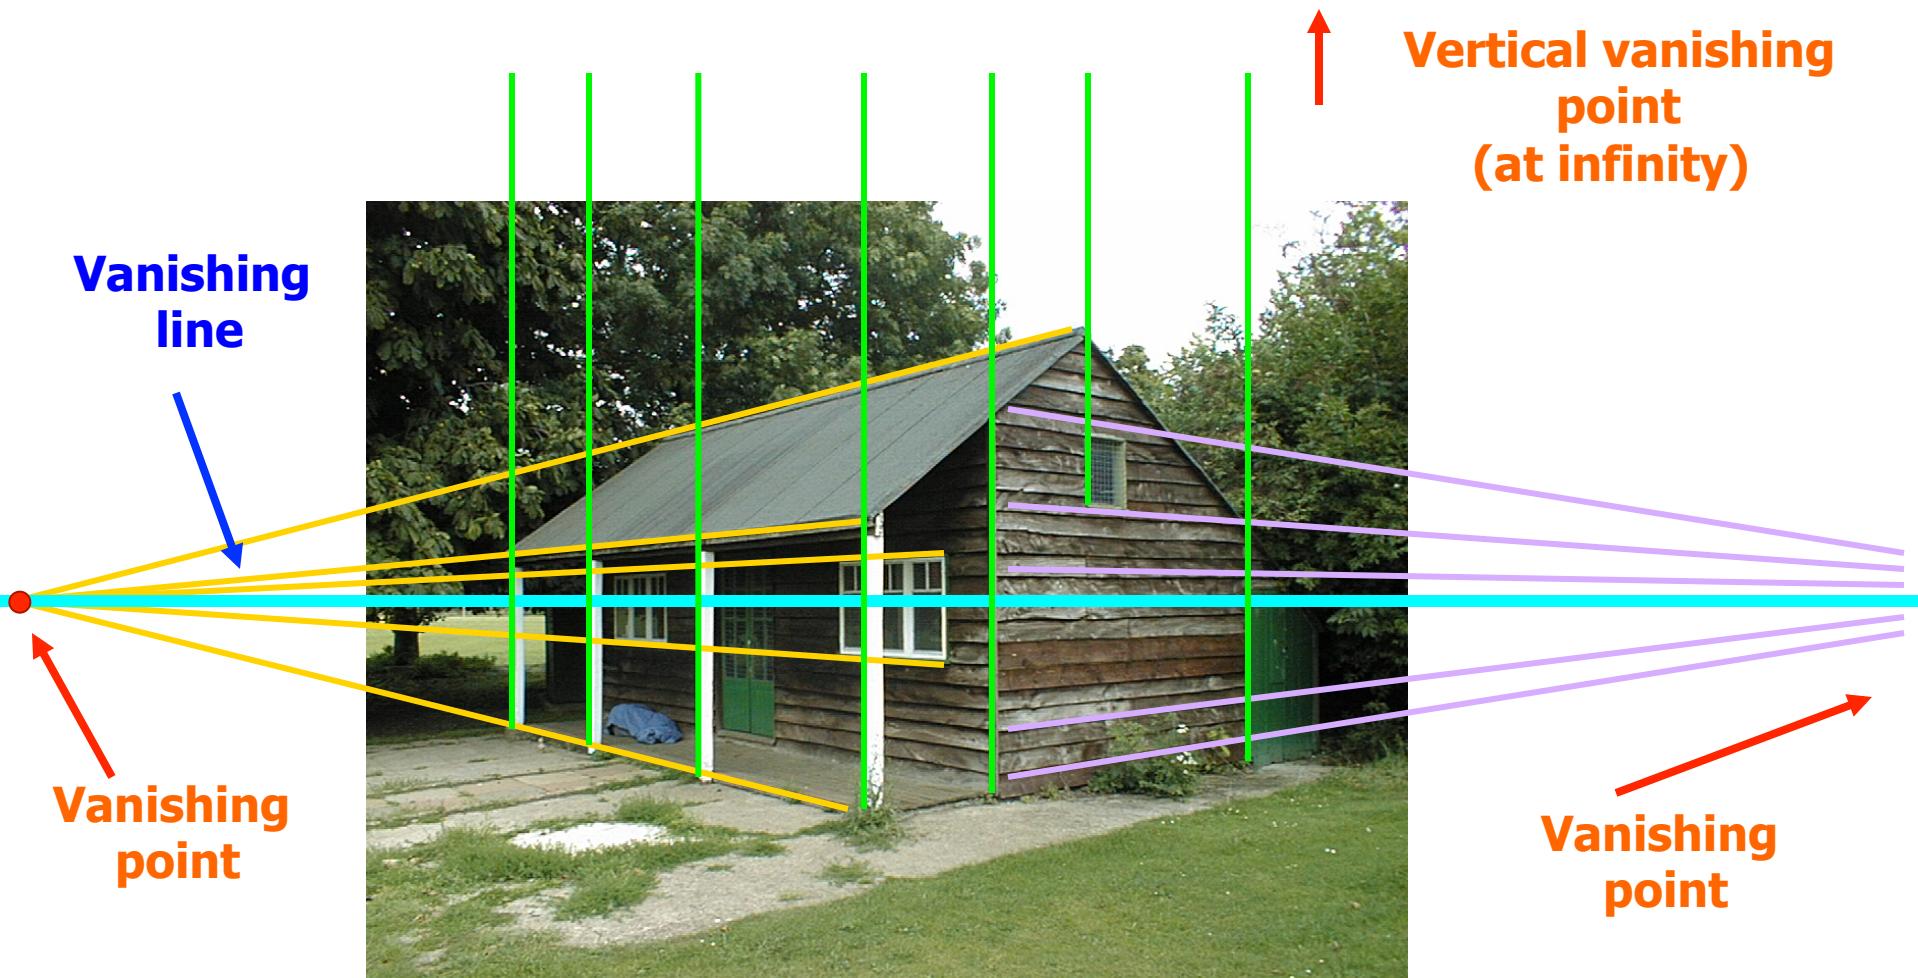
Vanishing points and lines

Parallel lines in the world intersect in the image at a “vanishing point”




Slide credit: J. Hayes

Vanishing points and lines



Slide from Efros, Photo from Criminisi

Pinhole camera

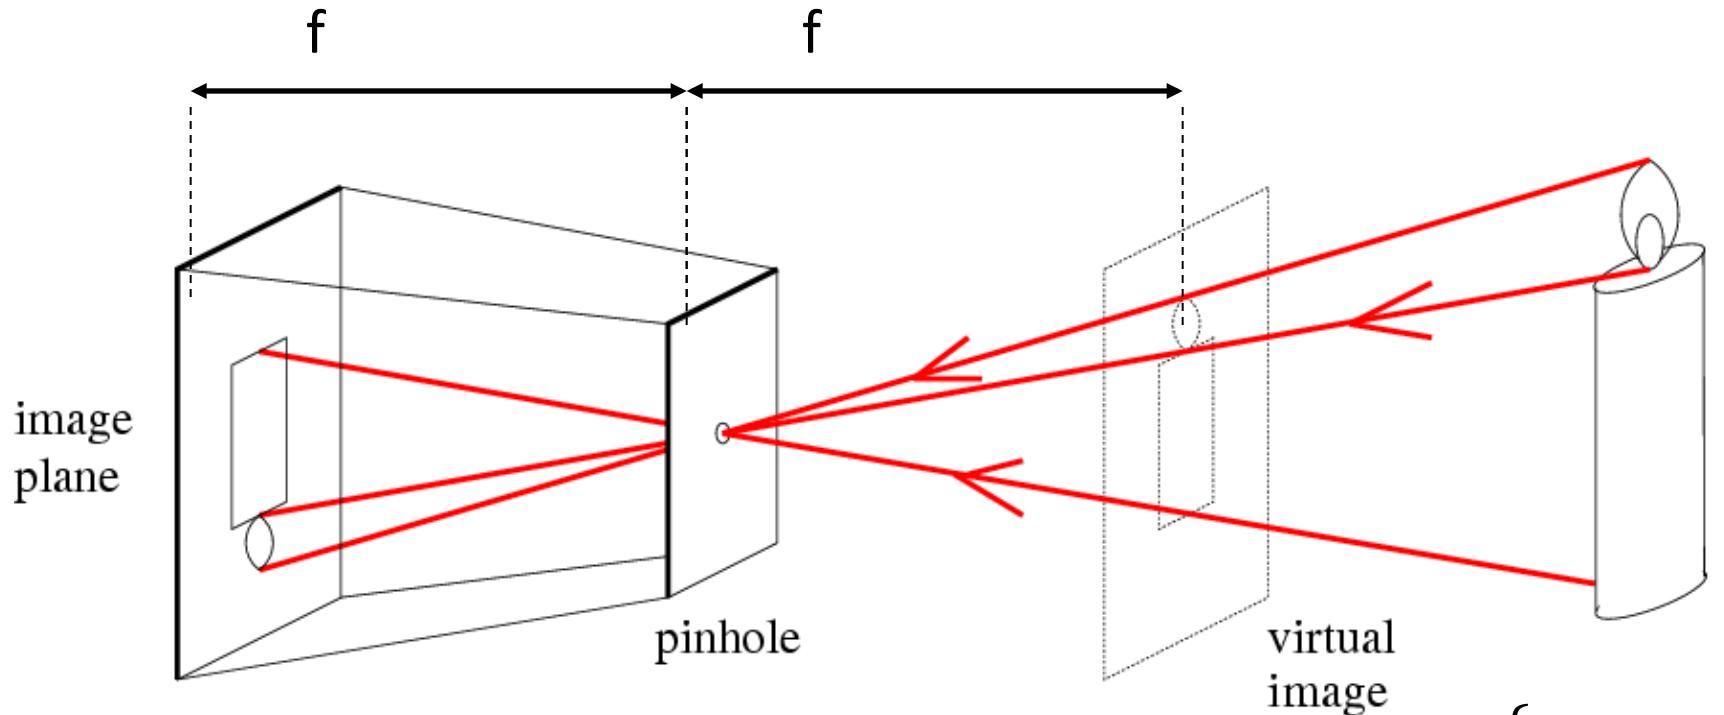


$$P = \begin{bmatrix} x \\ y \\ z \end{bmatrix} \rightarrow P' = \begin{bmatrix} x' \\ y' \end{bmatrix}$$
$$\left\{ \begin{array}{l} x' = f' \frac{x}{z} \\ y' = f' \frac{y}{z} \end{array} \right.$$

Note: z is always negative.

Derived using similar triangles

Pinhole camera

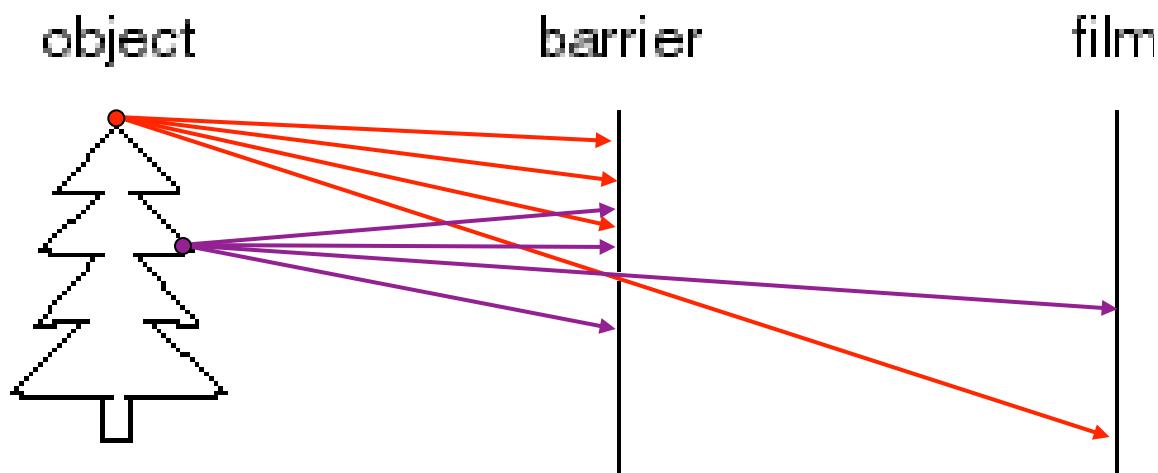


- Common to draw image plane *in front* of the focal point
- Moving the image plane merely scales the image.

$$\begin{cases} x' = f \frac{x}{z} \\ y' = f \frac{y}{z} \end{cases}$$

Pinhole camera

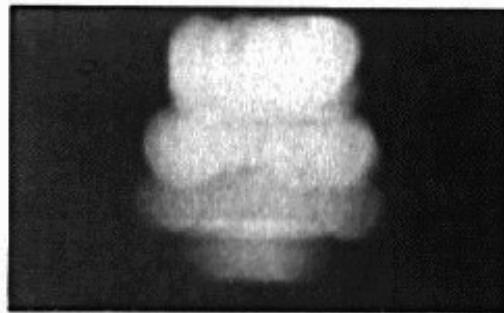
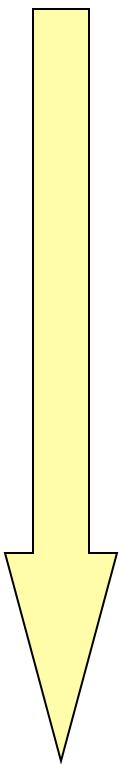
Is the size of the aperture important?



Cameras & Lenses

Shrinking
aperture
size

- Rays are mixed up



-Why the aperture cannot be too small?

- Less light passes through
- Diffraction effect

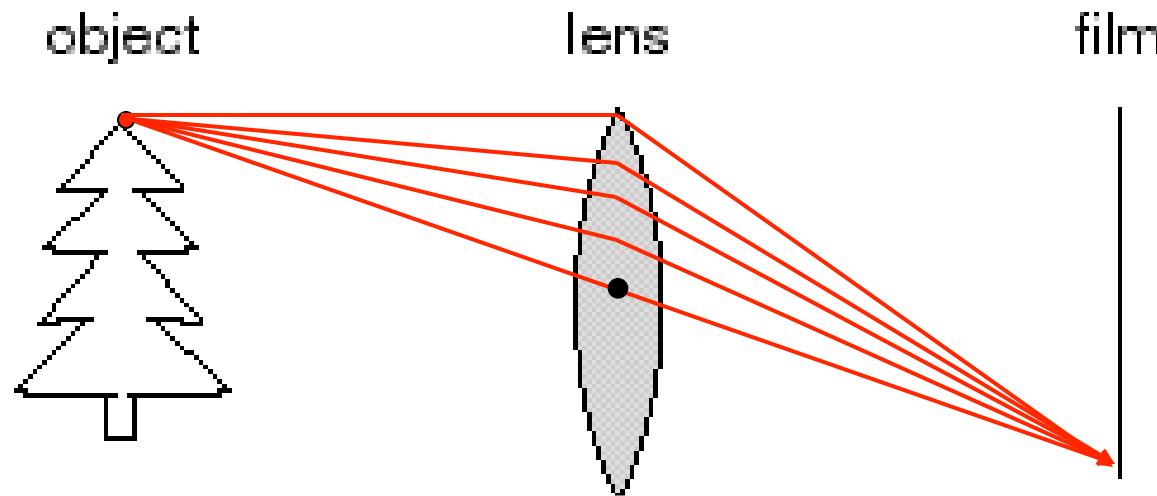
Adding lenses!

What we will learn today?

- Pinhole cameras
- Cameras & lenses
- The geometry of pinhole cameras
 - Projection matrix
 - Intrinsic parameters
 - Extrinsic parameters

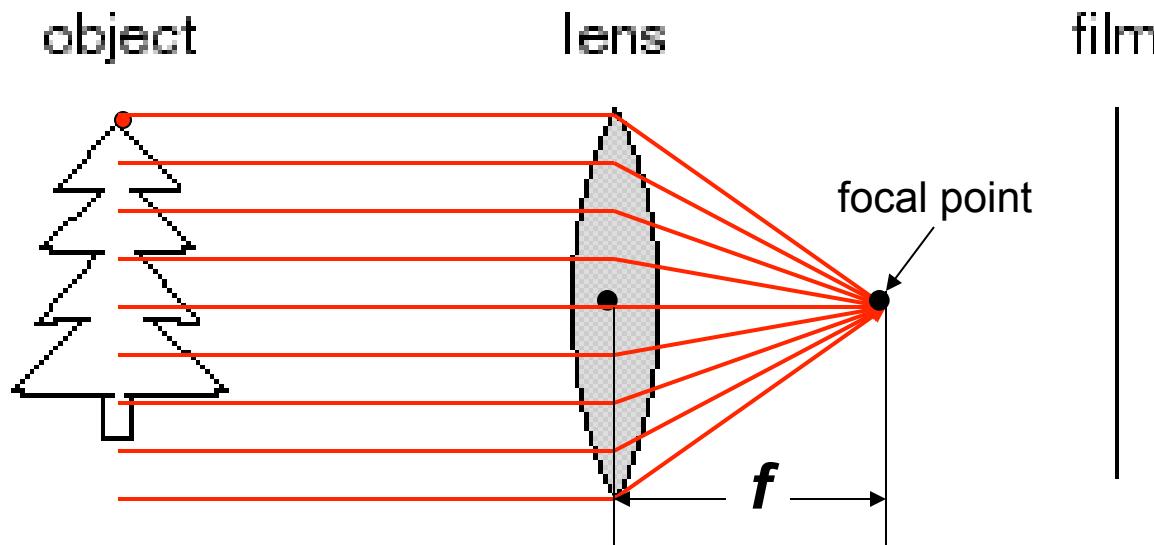
Reading:
[FP] Chapters 1 – 3
[HZ] Chapter 6

Cameras & Lenses



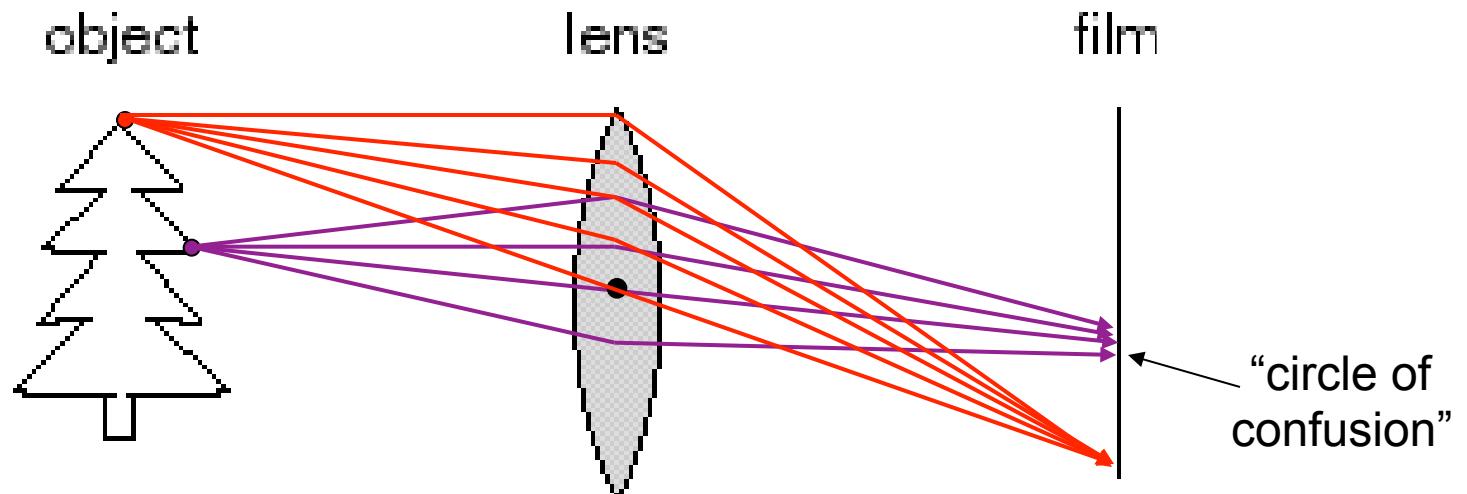
- A lens focuses light onto the film

Cameras & Lenses



- A lens focuses light onto the film
 - Rays passing through the center are not deviated
 - All parallel rays converge to one point on a plane located at the *focal length* f

Cameras & Lenses



- A lens focuses light onto the film
 - There is a specific distance at which objects are “in focus”
[other points project to a “circle of confusion” in the image]

Cameras & Lenses

- Laws of geometric optics
 - Light travels in straight lines in homogeneous medium
 - Reflection upon a surface: incoming ray, surface normal, and reflection are co-planar
 - Refraction: when a ray passes from one medium to another

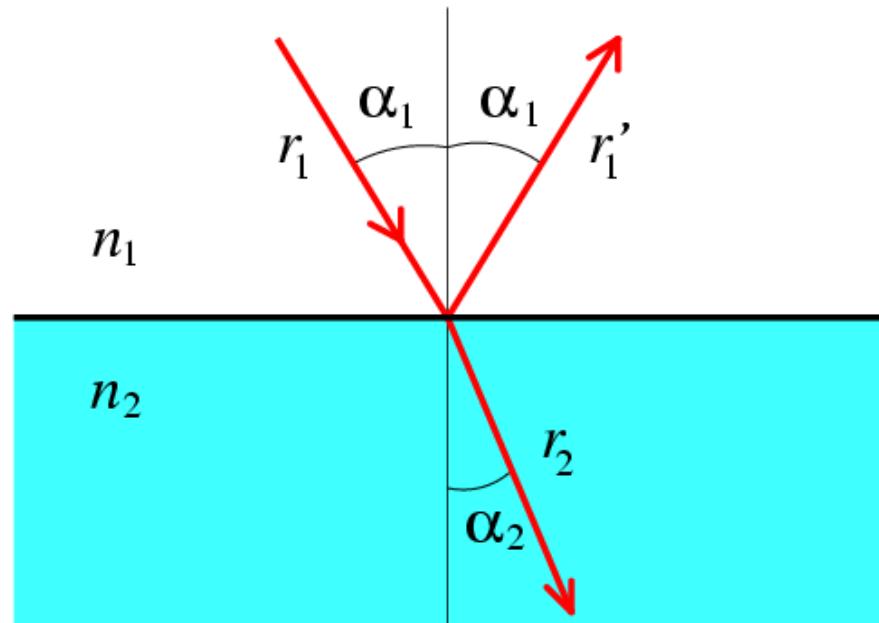
Snell's law

$$n_1 \sin \alpha_1 = n_2 \sin \alpha_2$$

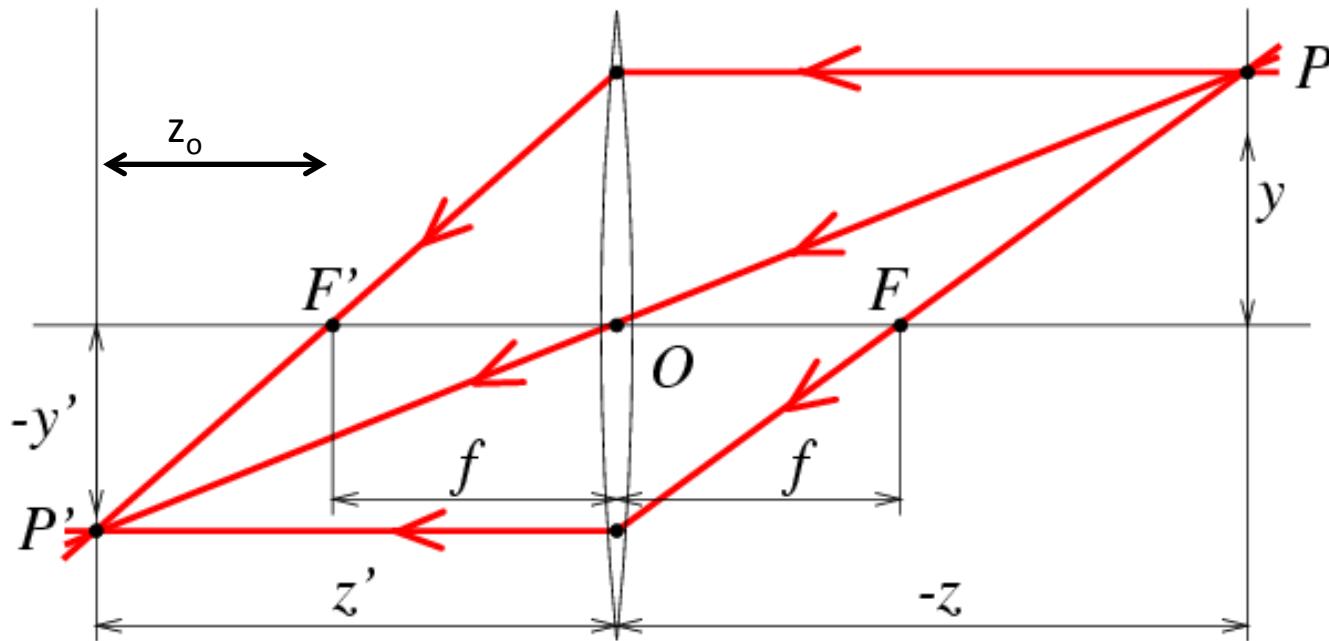
α_1 = incident angle

α_2 = refraction angle

n_i = index of refraction



Thin Lenses



$$z' = f + z_0$$

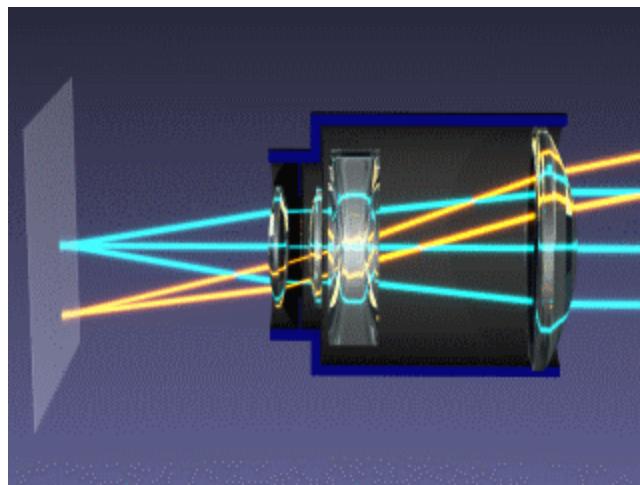
$$f = \frac{R}{2(n - 1)}$$

Snell's law:

$$n_1 \sin \alpha_1 = n_2 \sin \alpha_2$$

$$\begin{cases} \text{Small angles:} \\ n_1 \alpha_1 \approx n_2 \alpha_2 \\ n_1 = n \text{ (lens)} \\ n_1 = 1 \text{ (air)} \end{cases} \rightarrow \begin{cases} x' = z' \frac{x}{z} \\ y' = z' \frac{y}{z} \end{cases}$$

Cameras & Lenses

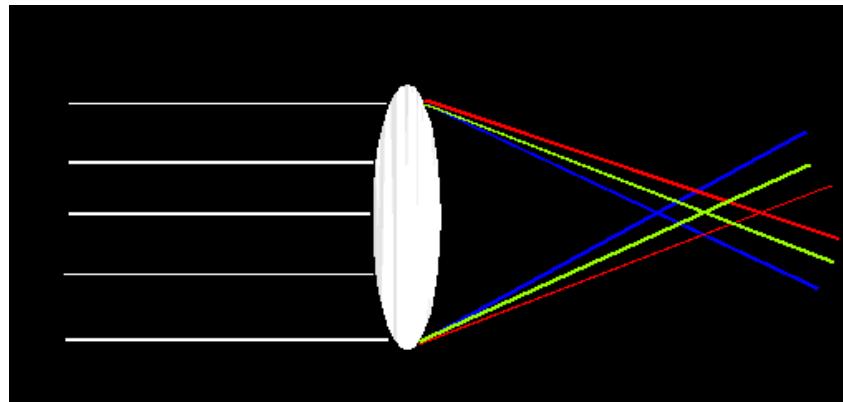


Source wikipedia

Issues with lenses: Chromatic Aberration


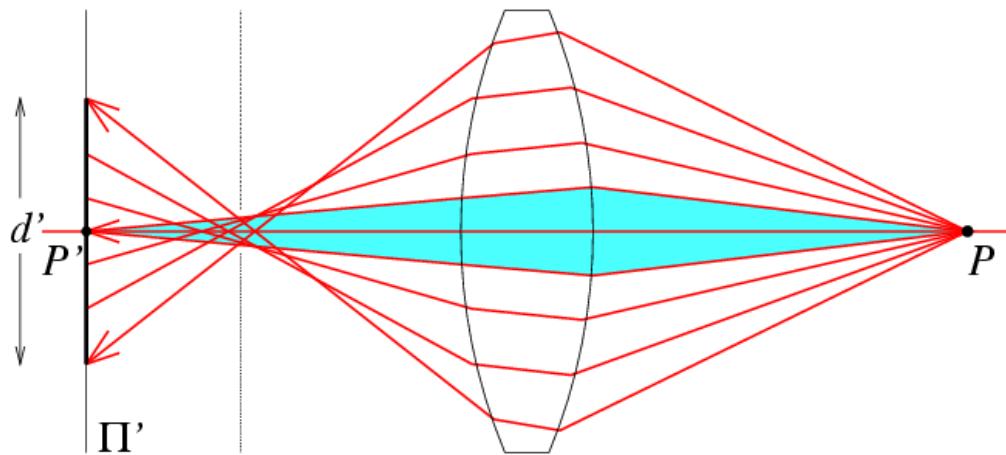
- Lens has different refractive indices for different wavelengths: causes color fringing

$$f = \frac{R}{2(n - 1)}$$



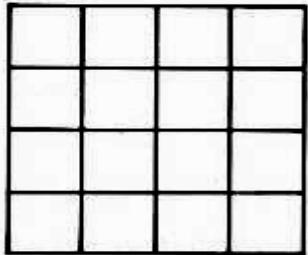
Issues with lenses: Chromatic Aberration

- Rays farther from the optical axis focus closer

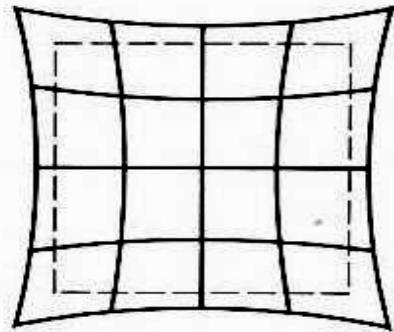


Issues with lenses: Chromatic Aberration

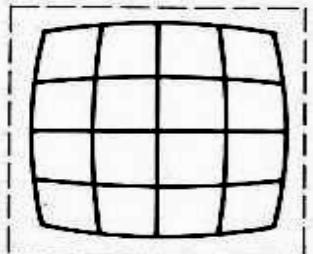
- Deviations are most noticeable for rays that pass through the edge of the lens



No distortion



Pin cushion



Barrel (fisheye lens)

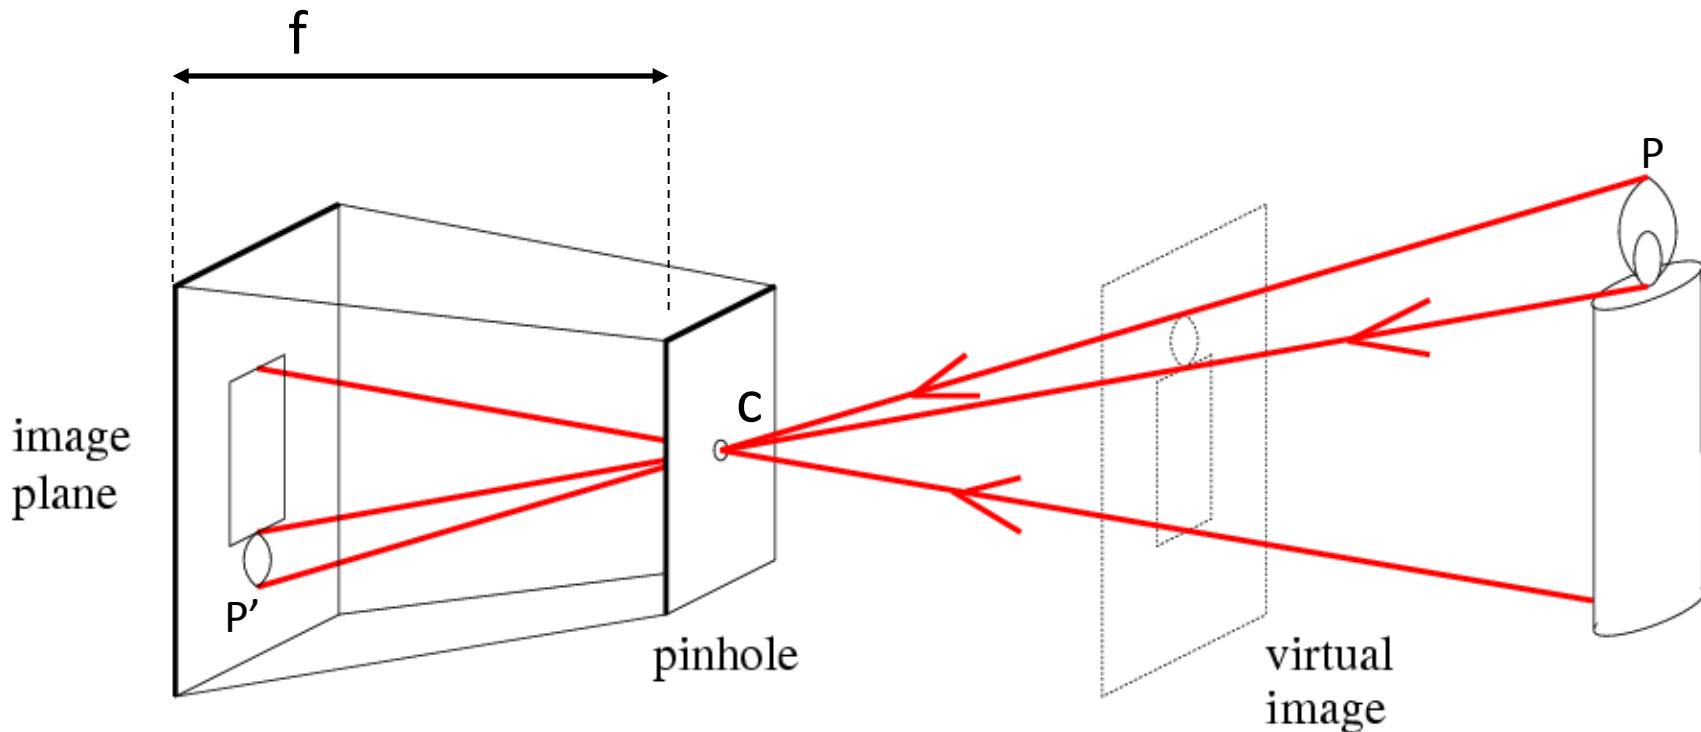
Image magnification decreases with distance from the optical axis



What we will learn today?

- Pinhole cameras
- Cameras & lenses
- The geometry of pinhole cameras
 - Projection matrix
 - Intrinsic parameters
 - Extrinsic parameters

Relating real-world point to a point on a camera



$$P = (x, y, z) \rightarrow P' = \left(f \frac{x}{z}, f \frac{y}{z} \right)$$

f = focal length

c = center of the camera

$$\mathfrak{R}^3 \xrightarrow{E} \mathfrak{R}^2$$

Relating real-world point to a point on a camera

Is this a linear transformation?

$$P = (x, y, z) \rightarrow P' = (f \frac{x}{z}, f \frac{y}{z})$$

No — division by z is nonlinear!

How to make it linear?

Homogeneous coordinates – a reminder

$$(x, y) \Rightarrow \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

homogeneous image
coordinates

$$(x, y, z) \Rightarrow \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$

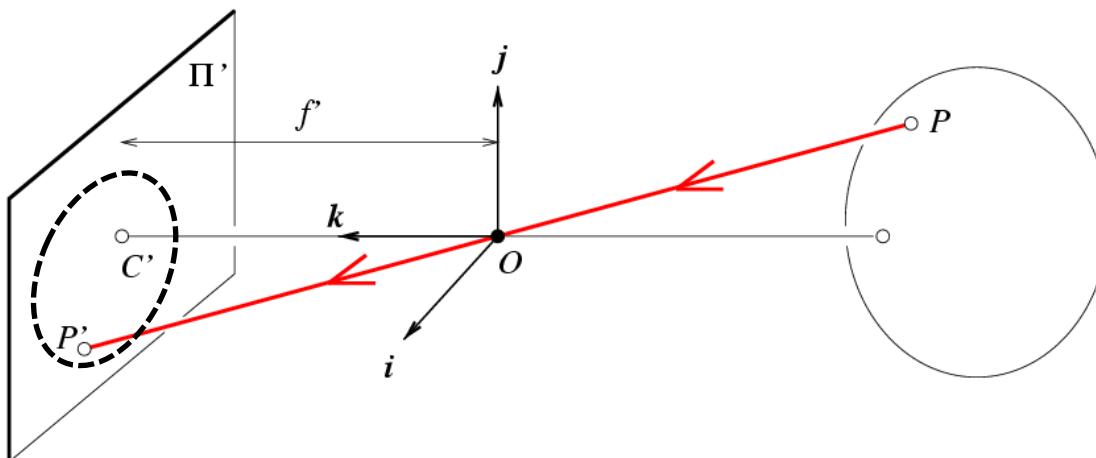
homogeneous scene
coordinates

- Converting *from* homogeneous coordinates

$$\begin{bmatrix} x \\ y \\ w \end{bmatrix} \Rightarrow (x/w, y/w)$$

$$\begin{bmatrix} x \\ y \\ z \\ w \end{bmatrix} \Rightarrow (x/w, y/w, z/w)$$

Relating a real-world point to a point on the camera



In Cartesian coordinates:

$$P = (x, y, z) \rightarrow P' = \left(f \frac{x}{z}, f \frac{y}{z} \right)$$

In homogeneous coordinates:

$$P' = \begin{bmatrix} f & x \\ f & y \\ z \end{bmatrix} = \underbrace{\begin{bmatrix} f & 0 & 0 & 0 \\ 0 & f & 0 & 0 \\ 0 & 0 & 1 & 0 \end{bmatrix}}_M \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$

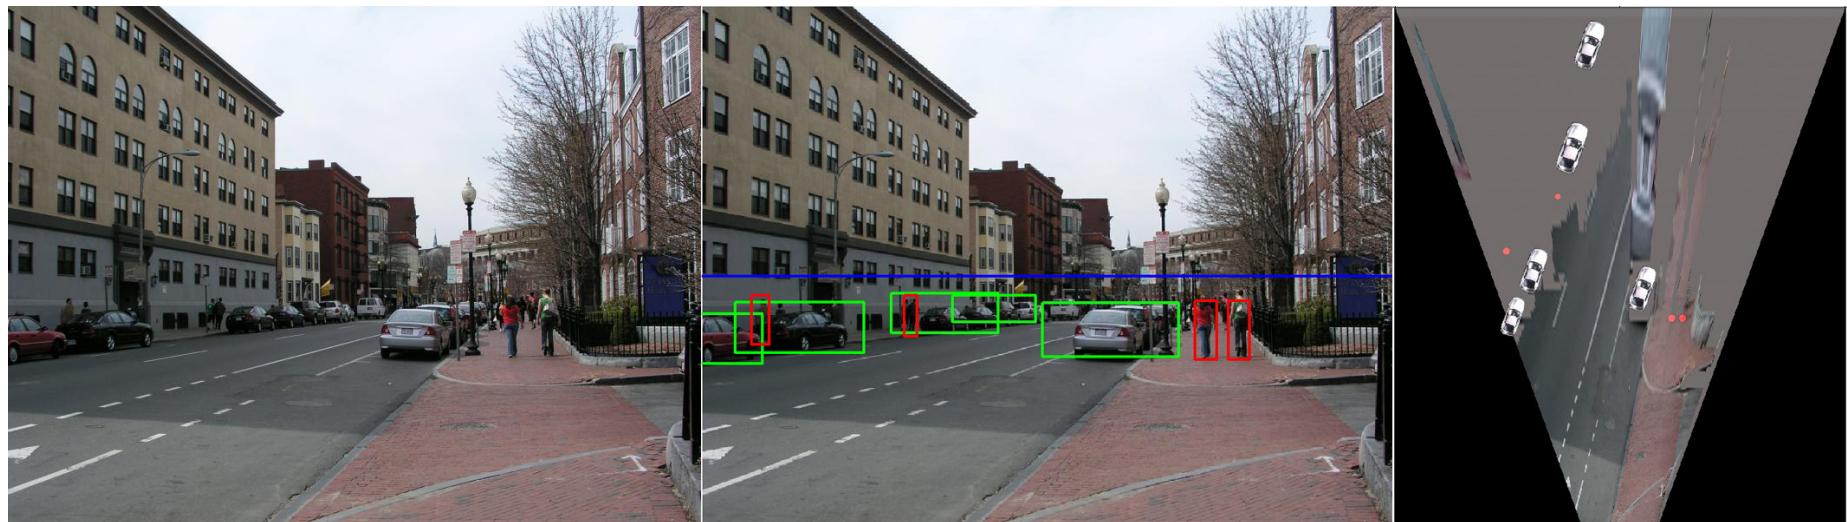
“Projection matrix”

$$\downarrow$$
$$P' = M P$$

$$\Re^4 \xrightarrow{H} \Re^3$$

Interlude: why does this matter?

Object Recognition (CVPR 2006)



Slide credit: J. Hayes

Inserting photographed objects into images (SIGGRAPH 2007)




Original



Created

Slide credit: J. Hayes

Relating a real-world point to a point on the camera



In homogeneous coordinates:

$$P' = \begin{bmatrix} f & x \\ f & y \\ z \end{bmatrix} = \underbrace{\begin{bmatrix} f & 0 & 0 & 0 \\ 0 & f & 0 & 0 \\ 0 & 0 & 1 & 0 \end{bmatrix}}_M \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$

ideal world

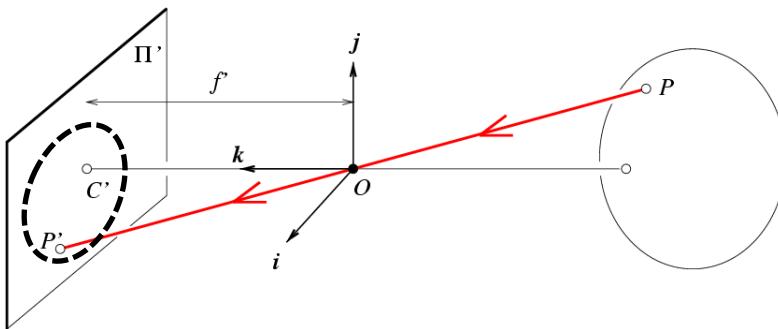
Intrinsic Assumptions

- Unit aspect ratio
- Optical center at (0,0)
- No skew

Extrinsic Assumptions

- No rotation
- Camera at (0,0,0)

Relating a real-world point to a point on the camera



In homogeneous coordinates:

$$P' = \begin{bmatrix} f & x \\ f & y \\ z \end{bmatrix} = \begin{bmatrix} f & 0 & 0 \\ 0 & f & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} 0 \\ 0 \\ 0 \\ 1 \end{bmatrix} = K \begin{bmatrix} I & 0 \end{bmatrix} P$$

K

Intrinsic Assumptions

- Unit aspect ratio
- Optical center at (0,0)
- No skew

Extrinsic Assumptions

- No rotation
- Camera at (0,0,0)

Remove assumption: known optical center

Intrinsic Assumptions

- Optical center at $(0,0)$
- Optical center at (u_0, v_0)
- Square pixels
- No skew

Extrinsic Assumptions

- No rotation
- Camera at $(0,0,0)$

$$P' = K \begin{bmatrix} I & 0 \end{bmatrix} P \xrightarrow{\text{blue arrow}} w \begin{bmatrix} u \\ v \\ 1 \end{bmatrix} = \begin{bmatrix} f & 0 & u_0 \\ 0 & f & v_0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$

Slide inspiration: S. Savarese

Remove assumption: square pixels

Intrinsic Assumptions

- Optical center at (u_0, v_0)
- ~~Square pixels~~
- **Rectangular pixels**
- No skew

Extrinsic Assumptions

- No rotation
- Camera at $(0,0,0)$

$$P' = K \begin{bmatrix} I & 0 \end{bmatrix} P \xrightarrow{\text{blue arrow}} w \begin{bmatrix} u \\ v \\ 1 \end{bmatrix} = \begin{bmatrix} \alpha & 0 & u_0 \\ 0 & \beta & v_0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$

The diagram shows the camera projection matrix P' as a product of the intrinsic matrix K and the extrinsic matrix P . The resulting equation is $w \begin{bmatrix} u \\ v \\ 1 \end{bmatrix} = \begin{bmatrix} \alpha & 0 & u_0 \\ 0 & \beta & v_0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$. The matrix $\begin{bmatrix} \alpha & 0 & u_0 \\ 0 & \beta & v_0 \\ 0 & 0 & 1 \end{bmatrix}$ is highlighted with a red dashed box. The parameters α and β are circled in red.

Slide inspiration: S. Savarese

Remove assumption: non-skewed pixels

Intrinsic Assumptions

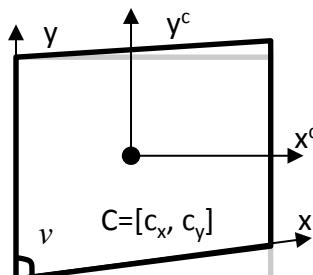
- Optical center at (u_0, v_0)
- Rectangular pixels
- ~~No skew~~
- Small skew

Extrinsic Assumptions

- No rotation
- Camera at $(0,0,0)$

$$P' = K \begin{bmatrix} I & 0 \end{bmatrix} P \rightarrow$$

$$w \begin{bmatrix} u \\ v \\ 1 \end{bmatrix} = \begin{bmatrix} \alpha & s & u_0 \\ 0 & \beta & v_0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$



Slide inspiration: S. Savarese

Remove assumption: non-skewed pixels

Intrinsic Assumptions

- Optical center at (u_0, v_0)
- Rectangular pixels
- Small skew

Extrinsic Assumptions

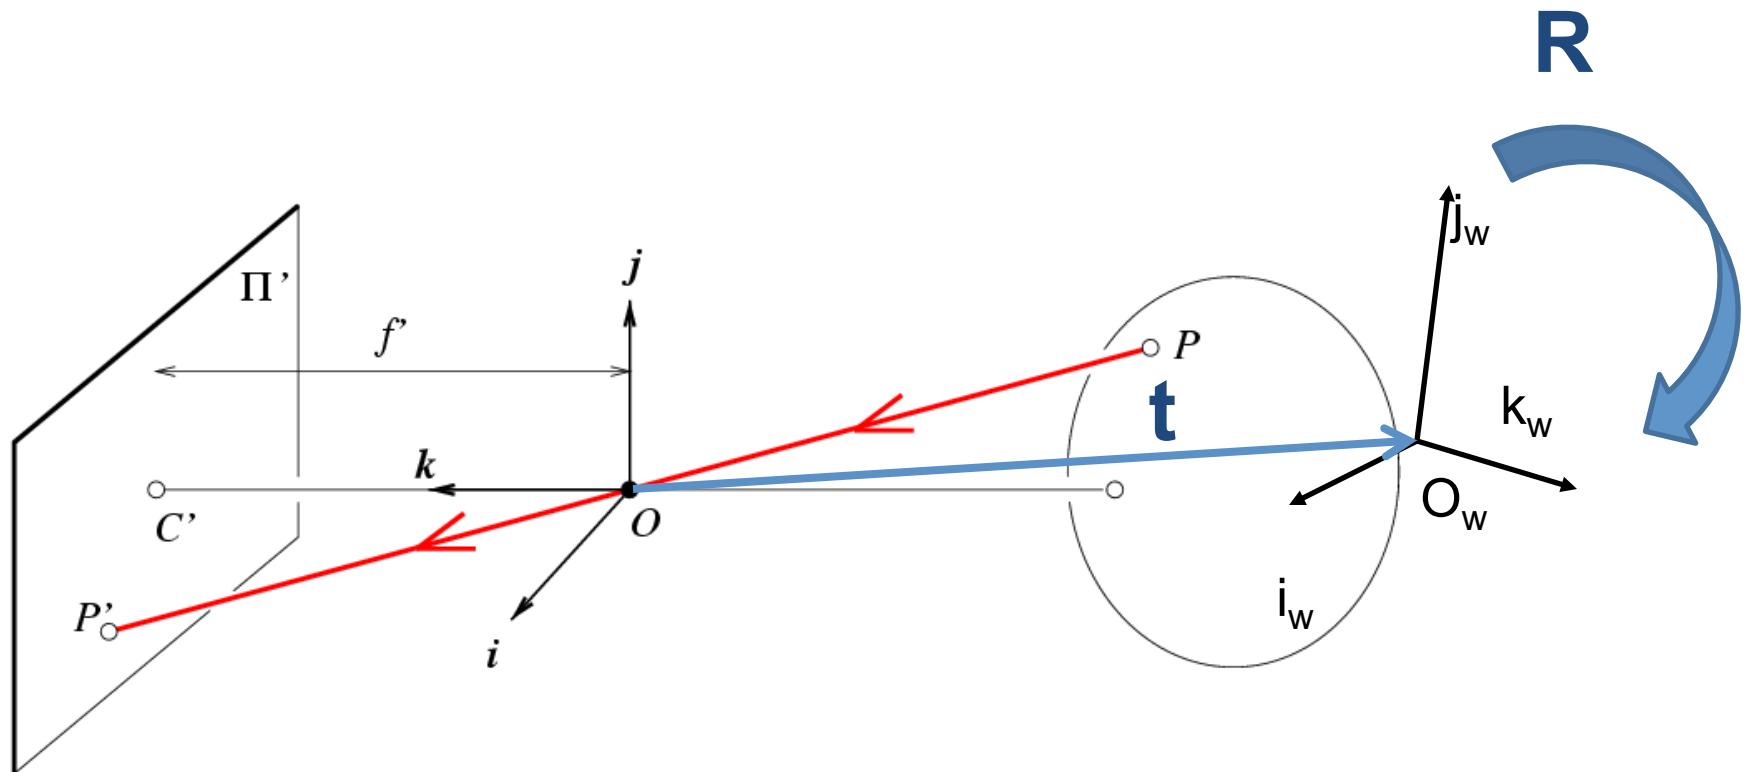
- No rotation
- Camera at $(0,0,0)$

$P' = K \begin{bmatrix} I & 0 \end{bmatrix} P \rightarrow w \begin{bmatrix} u \\ v \\ 1 \end{bmatrix} = \begin{bmatrix} \alpha & s & u_0 \\ 0 & \beta & v_0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$

Intrinsic parameters

Slide inspiration: S. Savarese

Real world camera: Translate + Rotate



Slide inspiration: S. Savarese, J. Hayes

Remove assumption: allow translation

Intrinsic Assumptions

- Optical center at (u_0, v_0)
- Rectangular pixels
- Small skew

Extrinsic Assumptions

- No rotation
- Camera at $(0,0,0)$ $\rightarrow (t_x, t_y, t_z)$


$$P' = K \begin{bmatrix} I & \bar{t} \end{bmatrix} P \quad \Rightarrow \quad w \begin{bmatrix} u \\ v \\ 1 \end{bmatrix} = \begin{bmatrix} \alpha & 0 & u_0 \\ 0 & \beta & v_0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ t_x \\ y \\ t_y \\ z \\ t_z \\ 1 \end{bmatrix}$$

Slide inspiration: S. Savarese

Remove assumption: allow rotation

Intrinsic Assumptions

- Optical center at (u_0, v_0)
- Rectangular pixels
- Small skew



Rotation around the coordinate axes, **counter-clockwise**

Extrinsic Assumptions

- ~~No~~ rotation
- Camera at (t_x, t_y, t_z)

$$R_x(\alpha) = \begin{bmatrix} 1 & 0 & 0 \\ 0 & \cos \alpha & -\sin \alpha \\ 0 & \sin \alpha & \cos \alpha \end{bmatrix}$$

$$R_y(\beta) = \begin{bmatrix} \cos \beta & 0 & \sin \beta \\ 0 & 1 & 0 \\ -\sin \beta & 0 & \cos \beta \end{bmatrix}$$

$$R_z(\gamma) = \begin{bmatrix} \cos \gamma & -\sin \gamma & 0 \\ \sin \gamma & \cos \gamma & 0 \\ 0 & 0 & 1 \end{bmatrix}$$

Slide inspiration: S. Savarese

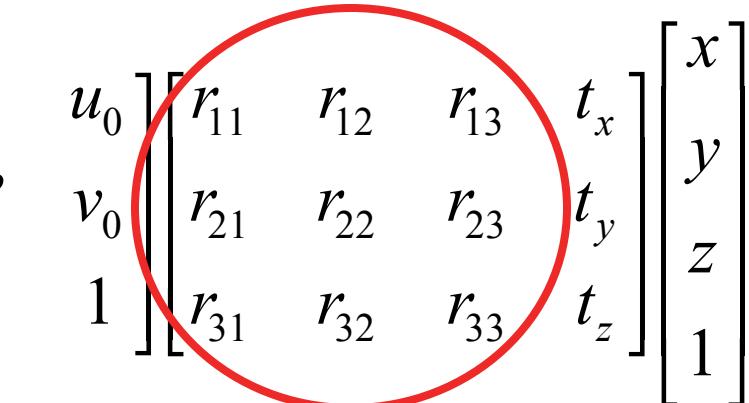
Remove assumption: allow rotation

Intrinsic Assumptions

- Optical center at (u_0, v_0)
- Rectangular pixels
- Small skew

Extrinsic Assumptions

- ~~No~~ rotation
- Camera at (t_x, t_y, t_z)

$$P' = K \begin{bmatrix} R & \bar{t} \end{bmatrix} P \rightarrow w \begin{bmatrix} u \\ v \\ 1 \end{bmatrix} = \begin{bmatrix} \alpha & s & u_0 \\ 0 & \beta & v_0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} r_{11} & r_{12} & r_{13} \\ r_{21} & r_{22} & r_{23} \\ r_{31} & r_{32} & r_{33} \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$


Slide inspiration: S. Savarese

A generic projection matrix

Intrinsic Assumptions

- Optical center at (u_0, v_0)
- Rectangular pixels
- Small skew

Extrinsic Assumptions

- Allow rotation
- Camera at (t_x, t_y, t_z)

$$P' = K \begin{bmatrix} R & \bar{t} \end{bmatrix} P \xrightarrow{\text{blue arrow}} w \begin{bmatrix} u \\ v \\ 1 \end{bmatrix} = \begin{bmatrix} \alpha & s & u_0 \\ 0 & \beta & v_0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} r_{11} & r_{12} & r_{13} & t_x \\ r_{21} & r_{22} & r_{23} & t_y \\ r_{31} & r_{32} & r_{33} & t_z \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$

Slide inspiration: S. Savarese

A generic projection matrix

Intrinsic Assumptions

- Optical center at (u_0, v_0)
- Rectangular pixels
- Small skew

Extrinsic Assumptions

- Allow rotation
- Camera at (t_x, t_y, t_z)

$$P' = K \begin{bmatrix} R & \bar{t} \end{bmatrix} P \xrightarrow{\text{blue arrow}} w \begin{bmatrix} u \\ v \\ 1 \end{bmatrix} = \begin{bmatrix} \alpha & s & u_0 \\ 0 & \beta & v_0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} r_{11} & r_{12} & r_{13} & t_x \\ r_{21} & r_{22} & r_{23} & t_y \\ r_{31} & r_{32} & r_{33} & t_z \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$

Degrees of freedom??

Slide inspiration: S. Savarese

A generic projection matrix

Intrinsic Assumptions

- Optical center at (u_0, v_0)
- Rectangular pixels
- Small skew

Extrinsic Assumptions

- Allow rotation
- Camera at (t_x, t_y, t_z)

$$P' = K \begin{bmatrix} R & \bar{t} \end{bmatrix} P \rightarrow w \begin{bmatrix} u \\ v \\ 1 \end{bmatrix} = \begin{bmatrix} \alpha & s & u_0 \\ 0 & \textcolor{red}{\beta} & v_0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} r_{11} & r_{12} & r_{13} & t_x \\ r_{21} & r_{22} & \textcolor{red}{r}_{23} & t_y \\ r_{31} & r_{32} & r_{33} & t_z \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$

Degrees of freedom??

Slide inspiration: S. Savarese

CS231a: Camera Calibration

estimate all intrinsic and extrinsic parameters

Intrinsic Assumptions

- Optical center at (u_0, v_0)
- Rectangular pixels
- Small skew

Extrinsic Assumptions


- Allow rotation
- Camera at (t_x, t_y, t_z)

$$P' = K \begin{bmatrix} R & \bar{t} \end{bmatrix} P \rightarrow w \begin{bmatrix} u \\ v \\ 1 \end{bmatrix} = \begin{bmatrix} \alpha & s & u_0 \\ 0 & \beta & v_0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} r_{11} & r_{12} & r_{13} & t_x \\ r_{21} & r_{22} & r_{23} & t_y \\ r_{31} & r_{32} & r_{33} & t_z \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$

Slide inspiration: S. Savarese

Orthographic Projection

- Special case of perspective projection
 - Distance from the COP to the image plane is infinite




- Also called “parallel projection”
- What’s the projection matrix?

$$w \begin{bmatrix} u \\ v \\ 1 \end{bmatrix} = \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$

Slide credit: Steve Seitz

Scaled Orthographic Projection

- Special case of perspective projection
 - Object dimensions are small compared to distance to camera

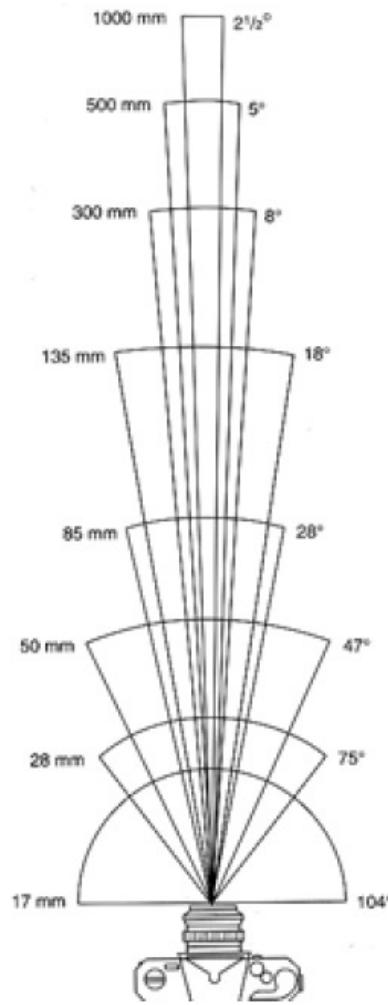


- Also called “weak perspective”
- What’s the projection matrix?

$$w \begin{bmatrix} u \\ v \\ 1 \end{bmatrix} = \begin{bmatrix} f & 0 & 0 & 0 \\ 0 & f & 0 & 0 \\ 0 & 0 & 0 & s \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$

Slide credit: Steve Seitz

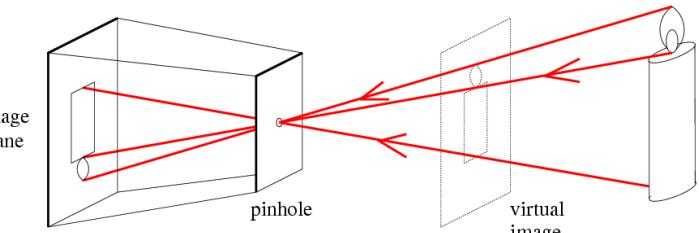

Field of View (Zoom)



From London and Upton

Things to remember

- Vanishing points and vanishing lines
- Pinhole camera model and camera projection matrix M
 - Intrinsic parameters
 - Extrinsic parameters
- Homogeneous coordinates



$$P' = K \begin{bmatrix} R & \bar{t} \end{bmatrix} P$$

$$(x, y) \Rightarrow \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

Slide inspiration: J. Hayes

What we have learned today?

- Pinhole cameras
- Cameras & lenses
- The geometry of pinhole cameras
 - Projection matrix
 - Intrinsic parameters
 - Extrinsic parameters

Reading:
[FP] Chapters 1 – 3
[HZ] Chapter 6